

Docker for Data Science

<http://bit.ly/d4ds-tutorial>

Aly Sivji



@CaiusSivjus

Joe Jasinski




@JoeJJasinski

Tathagata Dasgupta



@Tathagata



#d4ds



@



About Us

Aly Sivji @CaiusSivjus

Mathematician / Software Engineer @ divvyDOSE

Medical Informatics Grad Student @ Northwestern

Joe Jasinski @JoeJJasinski

Senior Software Engineer @ Nielsen

MS in Computer Science

Tathagata Dasgupta (aka T) @Tathagata

Senior Software Engineer @ Here Technologies

MS in Information Systems

Agenda

- Rules of the Road
- Data Science Overview
- Introduction to Docker
 - Hands-on Labs
- Data Science Workflows using Docker Containers
 - Hands-on Lab
- Break (3:00 pm)
- Docker Compose Overview
- PyCon Talk Recommender Application
 - Hands-on Labs

Rules of the Road

- Format: Lecture + Lab to reinforce concepts
- Main Github Repo: <http://bit.ly/d4ds-tutorial>
 - Setup instructions
 - Link to Slides: <http://bit.ly/d4ds-slides>
- Asking for Help
 - Raise your hand during lab sessions
 - Question session at the end of (most) labs
- Beyond the scope
 - Specific questions about how to fit Docker into YOUR workflow
 - Let's discuss offline!

Docker for Data Science

build passing


Materials for "Docker for Data Science" tutorial presented at PyCon 2018 in Cleveland, OH.

Slides

- [Description](#)
- [Audience](#)
- [Installation Instructions](#)
 - [Step 1: Install Docker and Docker-Compose](#)
 - [Step 2: Clone Git Repositories](#)
 - [Step 3: Download Docker Images](#)

Data Science

What is Data Science?



Source: [Drew Conway](#)

Data Science Use Cases



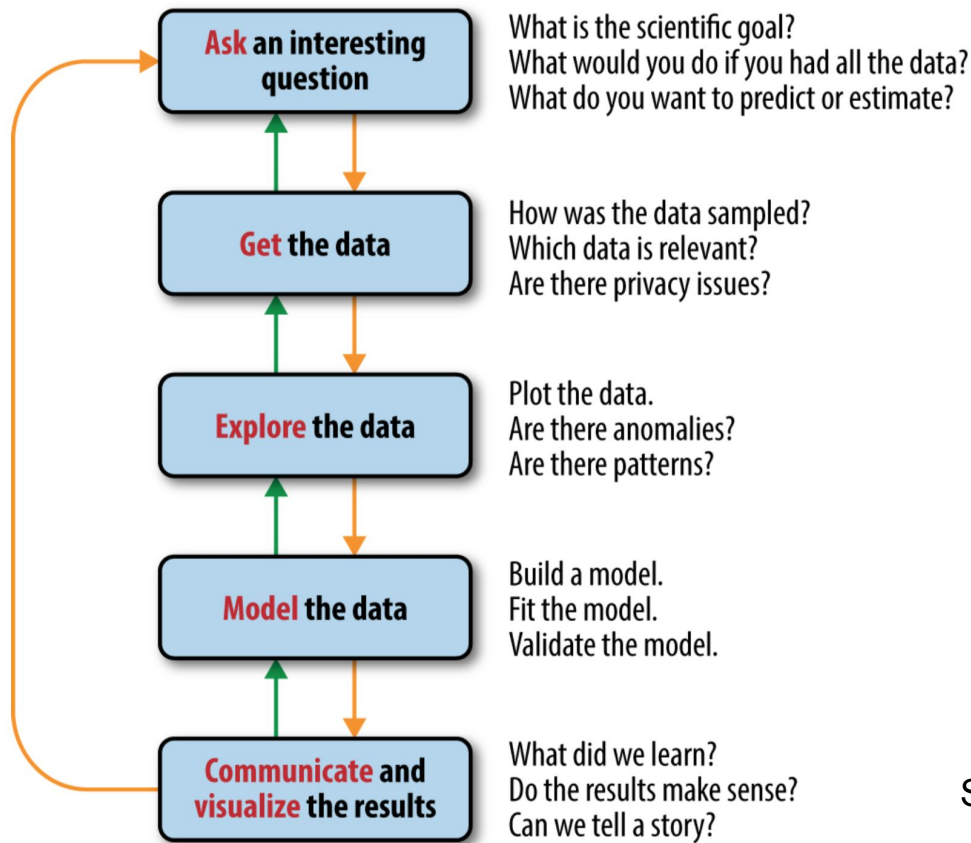
Data Science is Science

- Have a question
- Output is **findings + methodology**
- Reproducibility matters

Data Science Reproducibility



- Communicate results
- Defend decision making
- Auditable workflow

Data Science Process



Source: [Harvard CS 109](#)

Data Science and Python

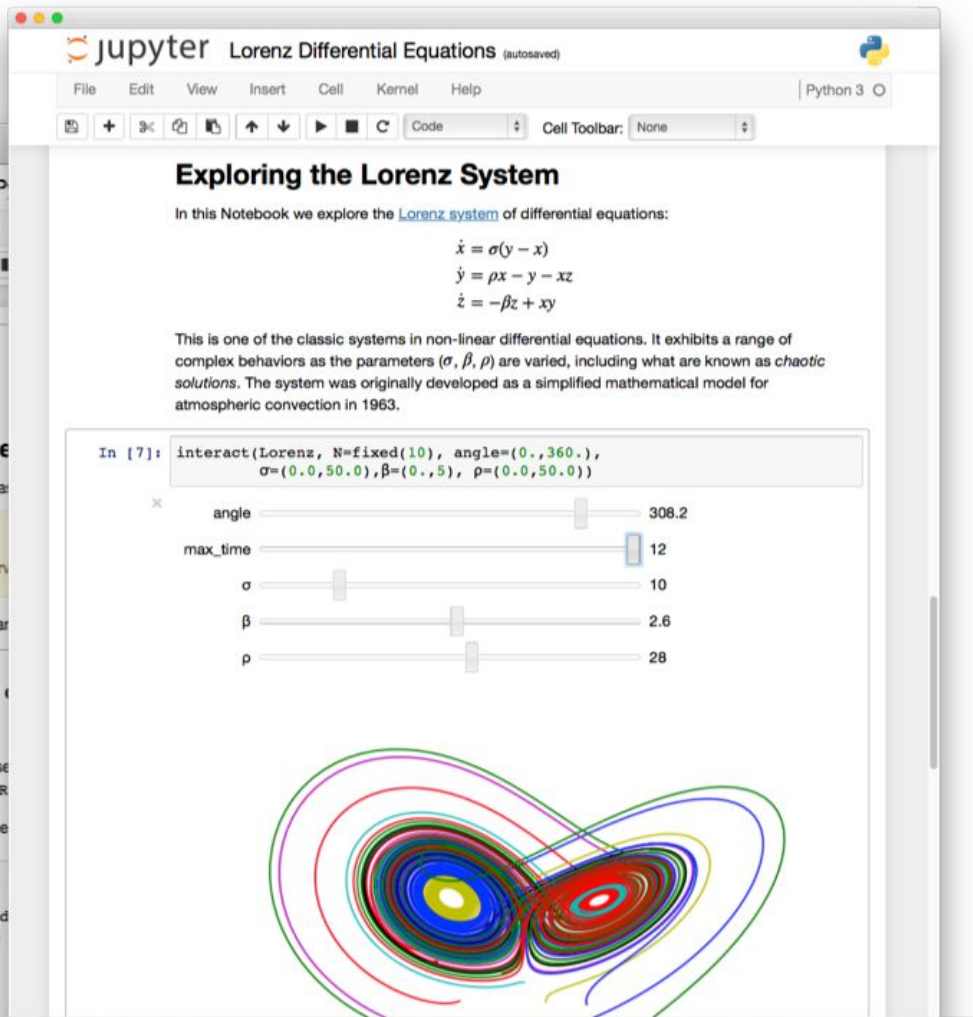
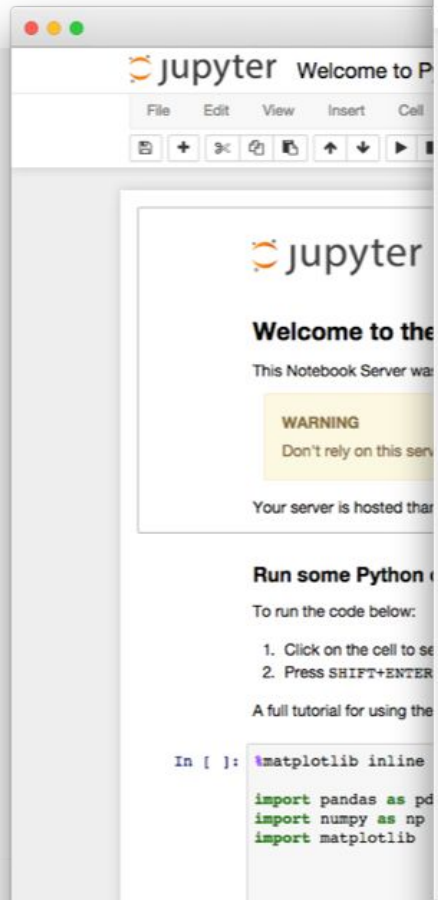


IP[y]:
IPython




Jupyter Notebooks

- Create / Share documents containing:
 - Live code
 - Equations
 - Visualizations
 - Explanatory Text
- Perfect for Data Science Workflows



Jupyter Limitations




Docker

Introduction to Docker

- Docker allows us to **package** and **run** applications in an **isolated environment**




Shipping Container Analogy




Source: [Docker](#)


Software Containers






Code



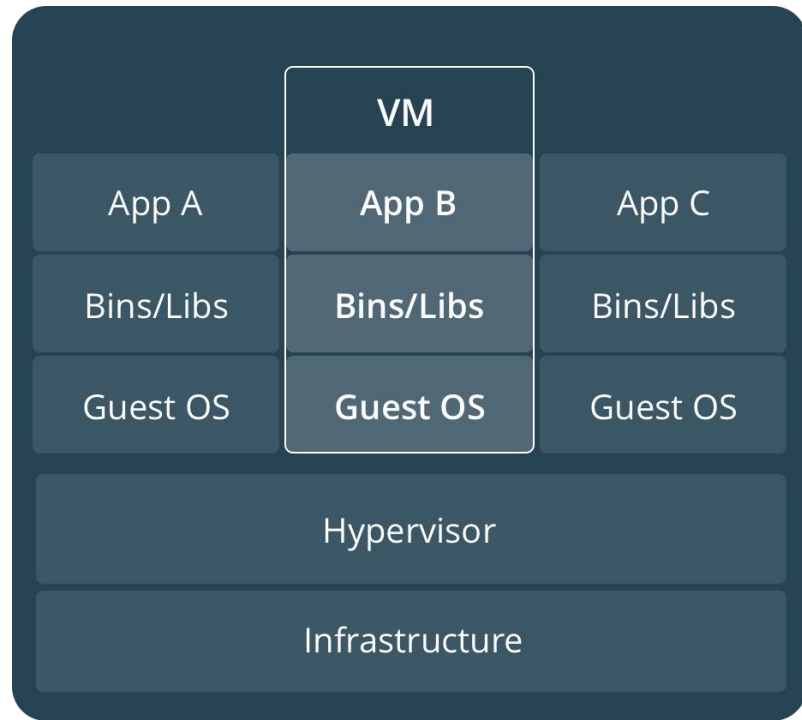
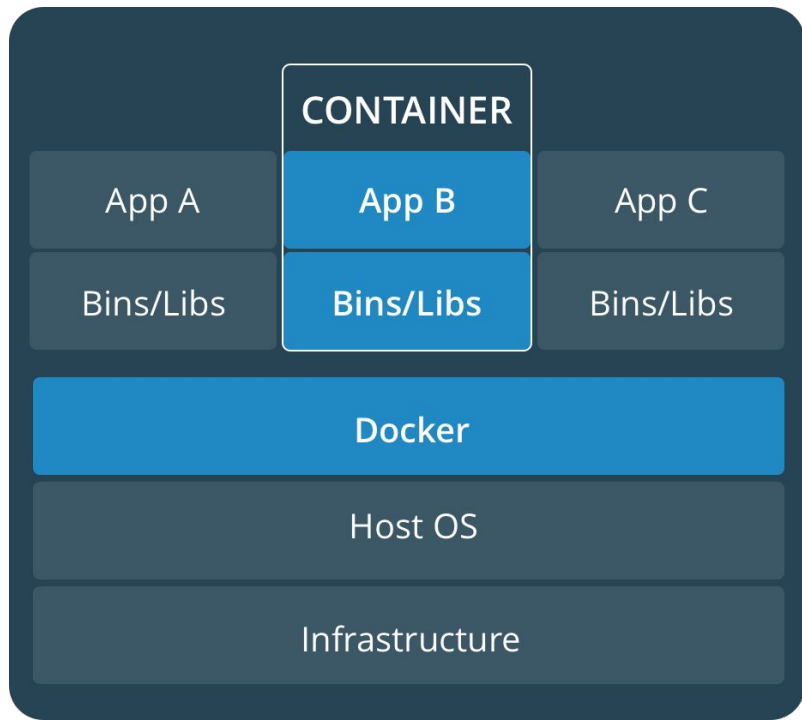
Dependencies



Data



Docker Containers vs Virtual Machines

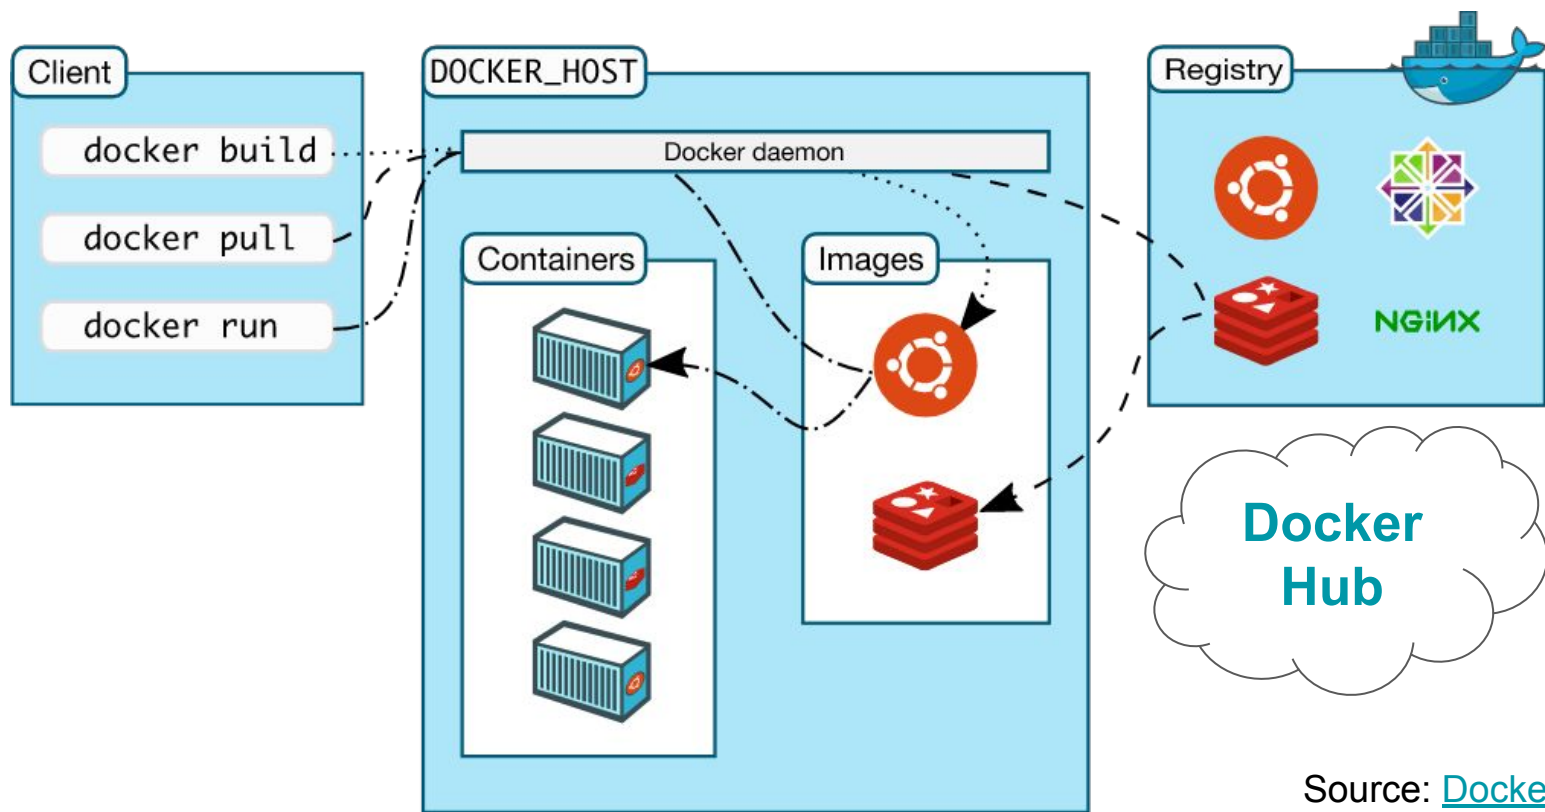


Source: [Docker Docs](#)

Docker Use Cases

- Streamline development workflows
 - Continuous Integration and Deployment (CI/CD)
- Microservices
 - But remember, monolith first
- Reproducible Data Science

Docker Architecture: Overview



Source: [Docker Docs](https://docs.docker.com/)



OFFICIAL REPOSITORY

python ☆

Last pushed: 5 hours ago

Repo Info

Tags

Short Description

Python is an interpreted, interactive, object-oriented, open-source programming language.

Full Description

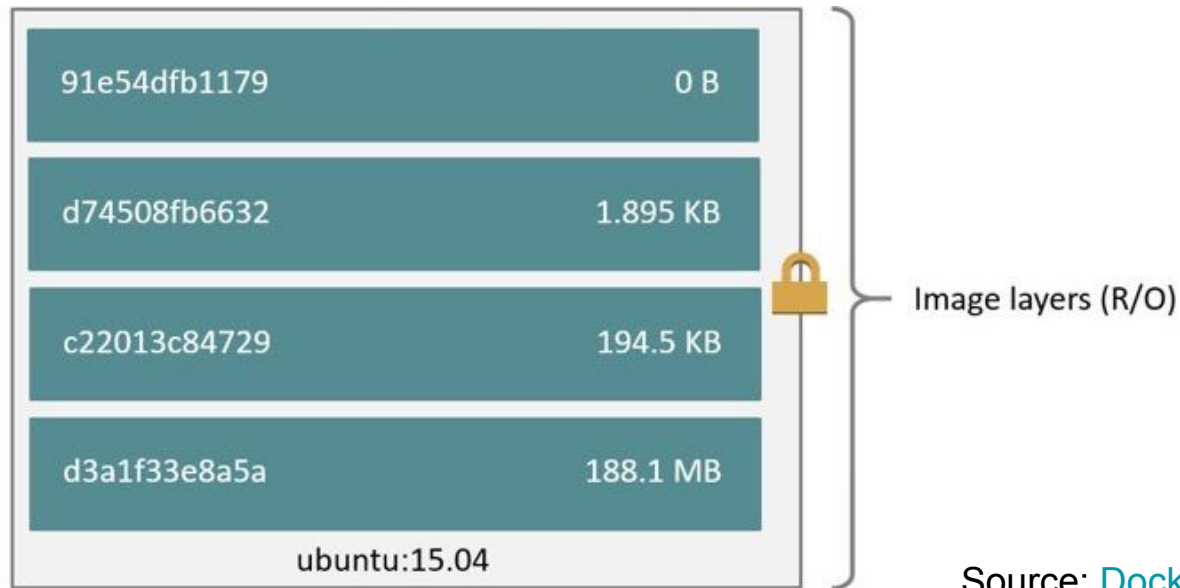
Supported tags and respective **Dockerfile** links

Simple Tags

- 3.7.0b3-stretch , 3.7-rc-stretch , rc-stretch ([3.7-rc/stretch/Dockerfile](#))
- 3.7.0b3-slim-stretch , 3.7-rc-slim-stretch , rc-slim-stretch , 3.7.0b3-slim , 3.7-rc-slim , rc-slim ([3.7-rc/stretch/slim/Dockerfile](#))
- 3.7.0b3-alpine3.7 , 3.7-rc-alpine3.7 , rc-alpine3.7 , 3.7.0b3-alpine , 3.7-rc-alpine , rc-alpine ([3.7-rc/alpine3.7/Dockerfile](#))

Docker Image


- A frozen snapshot of a container



Source: [Docker Docs](#)

Docker Containers

- Runtime instance: `docker run [image]`



Source: [Docker Docs](#)

Object-Oriented Programming Analogy

- Images : Classes
- Layers : Inheritance
- Containers : Objects

Lab: Docker Essentials

<http://bit.ly/d4ds-lab1>

Creating Docker Images

1. Freeze container using [docker commit](#)
2. [Dockerfile](#) and [docker build](#) * Preferred *
 - File containing all commands used to assemble image
 - Automated build

Dockerfile Commands

- FROM - sets base image
- LABEL - adds metadata to image
 - MAINTAINER is deprecated
 - LABEL maintainer="Aly Sivji <alysivji@gmail.com>"
- COPY - copies files / directories into image
 - .dockerignore
- ENV - sets environment variable
- WORKDIR - sets working directory

Dockerfile Commands

- RUN - executes shell commands in a new layer

```
RUN pip install jupyter
```

2 layers

```
RUN pip install pandas
```

```
RUN pip install jupyter && \  
    pip install pandas
```

1 layer

Dockerfile - Configuring Runtime

- [ENTRYPOINT](#) - configures container to run as executable
- [CMD](#) - provides default for executing container
 - [CMD and ENTRYPOINT interaction](#)
- Two forms:
 - Shell `CMD python hello-world.py`
 - Exec (preferred) `CMD ["python", "hello-world.py"]`
- [Additional Information](#)

Hello World Dockerfile

```
# Use latest Python runtime as base image
```

```
FROM python:3.6.3-alpine3.6
```

```
# Set the working directory to /app and copy current dir
```

```
WORKDIR /app
```

```
COPY . /app
```

```
# Run hello_world.py when the container launches
```

```
CMD ["python", "hello_world.py"]
```


Building Image

```
$ docker build -t hello-world .
```

```
Sending build context to Docker daemon  3.072kB
```

```
Step 1/4 : FROM python:3.6.3-alpine3.6
```

```
...
```

```
Successfully built f4e5a0ccfcd5
```

```
Successfully tagged hello-world:latest
```

Source: [Docker Docs](#)

Container Commands

- Create Container

```
$ docker run hello-world
```

```
Hello World!
```

- Restart Container

```
$ docker start -ia [CONTAINER]
```

```
$ docker run [OPTIONS] IMAGE [COMMAND]
```

- [Options]

-d	Detached (runs in background)
-a	Attach to STDOUT/STDERR
-i	Interactive (keeps STDIN open)
-t	Allocates pseudo-TTY
--name [NAME]	Set the container name

- [Command]

- Can pass in parameters or **/bin/sh** to get into container's shell

Managing Data Inside Containers

- Data disappears when we delete a container
- docker cp to copy files in/out of containers
- Mount data volume inside container

Adding Data Volume to Container

```
$ docker run -v /full/local/path:/mounted_dir
```

Host Path 

 Container Path

- Best Practice: Add VOLUME command to Dockerfile

```
# Create mount point for external volumes
```



```
VOLUME /mounted_dir
```

Binding Ports

- Setup port forwarding to connect to containers

```
$ docker run -p 9999:8888
```

Host Port



Container Port

- Best Practice: Add EXPOSE command to Dockerfile

```
# Make port 8888 available to outside world
```

```
EXPOSE 8888
```

Dockerfile - Best Practices

- Be explicit about build process
- Containers should be stateless
- Use `.dockerignore` file
- Avoid installing unnecessary packages
 - Clean cache after installation
- Each container should have only one concern / purpose
- Minimize the number of layers
 - Multi-line arguments, sort alphabetically
- CMD should be used to run processes inside container
 - Advanced users should use it in conjunction with ENTRYPOINT
- MAINTAINER is deprecated; use LABEL

Docker Container Lifecycle

Conception

BUILD an Image from a Dockerfile

Birth

RUN (create+start) a container

Reproduction

COMMIT (persist) a container to a new image

RUN a new container from an image

Sleep

KILL a running container

Wake

START a stopped container

Death

RM (delete) a stopped container

Extinction

RMI a container image (delete image)

Source: [Docker 101](#)

Docker Commands: Containers

Lifecycle

- ✧ [docker create](#)
- ✧ [docker rename](#)
- ✧ [docker run](#)
- ✧ [docker rm](#)
- [docker update](#)

Misc

- [docker cp](#)
- [docker export](#)
- ✧ [docker exec](#)

Start/Stop

- ✧ [docker start](#)
- [docker stop](#)
- [docker restart](#)
- [docker pause](#)
- [docker unpause](#)
- [docker wait](#)
- ✧ [docker kill](#)
- ✧ [docker attach](#)

Info

- ✧ [docker ps](#)
- [docker logs](#)
- ✧ [docker inspect](#)
- [docker events](#)
- [docker port](#)
- ✧ [docker top](#)
- [docker stats](#)
- [docker diff](#)

Source: [Docker Cheat Sheet](#)

Docker Commands: Images

Lifecycle

- ✧ [docker images](#)
- [docker import](#)
- ✧ [docker build](#)
- [docker commit](#)
- ✧ [docker rmi](#)
- [docker load](#)
- [docker save](#)

Info

- [docker history](#)
- [docker tag](#)

Registry

- [docker login](#)
- [docker logout](#)
- [docker search](#)
- ✧ [docker pull](#)
- ✧ [docker push](#)

Tips and Tricks


- Smaller images are better. Install only the packages you need.
 - Look into different Linux distributions ([Alpine Linux](#)... only 5MB!)
 - Clear cache after installing or use no-cache flags!
- [Link bash_history and keep track of commands typed inside container](#)
- [dockviz](#) command line app to visualize docker data
- Ctrl + P + Q to detach from container while inside shell
- [Instructions on mounting symbolic links](#)
- Always set IP address for apps running inside container to 0.0.0.0

Lab: Dockerfile Essentials


<http://bit.ly/d4ds-lab2>

Questions?

Data Science Workflows with Docker



+



Self-Contained Container (Workflow #1)

- Problem: Sharing results (Jupyter notebook)
- Workflow:
 - Create Docker image with libraries, data and notebook
 - Push image to DockerHub

Self-Contained Container: Dockerfile

```
FROM python:3.6.3-slim
```

```
LABEL maintainer="Aly Sivji <alysivji@gmail.com>"
```

```
WORKDIR /app
```

```
COPY . /app
```

```
RUN pip --no-cache-dir install numpy pandas seaborn sklearn jupyter
```

```
EXPOSE 8888
```

```
# Run app.py when the container launches
```

```
CMD ["jupyter", "notebook", "--ip='*'", "--port=8888",  
"--no-browser", "--allow-root"]
```


Self-Contained Container: Commands

- Build Image

```
$ docker build -t alysivji/workflow1-self-contained .
```

- Initialize Container

```
$ docker run -p 9999:8888  
alysivji/workflow1-self-contained
```

- Restart Container

```
$ docker start -ia [CONTAINER]
```

Self-Contained Container: Docker Hub

- Upload to Docker Hub

```
$ docker login
```

```
$ docker push [full-image-name]
```

- Download Image

```
$ docker pull [full-image-name]
```

- Instructions from previous slide for lifecycle

Data Science Project (Workflow #2)

- Problem:

- Need to standardize team development environment
- Project based workflows

- Workflow:

- Create team / project image with dev environment
- Mount volume containing notebooks and data

Data Science Project: Benefits

- Separate out projects
- Create container to onboard new employees
- Easy to upgrade dependencies
 - Build automated testing pipeline

Data Science Project: Dockerfile

```
FROM continuumio/miniconda3
```

```
LABEL maintainer="alysivji@gmail.com"
```

```
WORKDIR /app
```

```
RUN conda install jupyter -y && \  
    conda clean -y -all
```

```
EXPOSE 8888
```

```
VOLUME /app
```

```
CMD ["jupyter", "notebook", "--ip='*'", "--port=8888",  
    "--no-browser", "--allow-root"]
```

Data Science Project: Commands

- Build Image

```
$ docker build -t  
alysivji/workflow2-data-science-project .
```

- Initialize Container

```
$ docker run -p 9999:8888 -v  
/Users/alysivji/siv-dev/datasci:/app  
alysivji/workflow2-data-science-project
```

- Restart Container

```
$ docker start -ia [CONTAINER]
```

Data Driven App (Workflow #3)

- Problem: Distributing application
- Workflow:
 - Package app in image and deploy using Docker
- Further Reading
 - [Docker Compose](#)

Data Driven App: Dashboard

- Data stored on local machine
- Create & package dashboard inside container
 - [Dash Tutorial](#)
- Container is an executable on top of data
 - Start container to view dashboard

Data Driven App: Dockerfile

```
FROM python:3.6.3-alpine3.6
```

```
LABEL maintainer="alysivji@gmail.com"
```

```
WORKDIR /app
```

```
COPY . /app
```

```
RUN pip --no-cache-dir install -r /app/requirements.txt
```

```
EXPOSE 8050
```

```
VOLUME /app/data
```

```
ENTRYPOINT ["python"]
```

```
CMD ["plot_timeseries.py"]
```

Data Driven App: Commands

- Build Image

```
$ docker build -t alysivji/workflow3-data-driven-app .
```

- Initialize Container

```
$ docker run -p 8050:8050 -v  
/Users/alyshivji/siv-dev/docker-example:/app/data  
--name dashboard alysivji/workflow3-data-driven-app
```

- Restart Container

```
$ docker start -ia dashboard
```

Lab: Data Science Workflows using Docker

<http://bit.ly/d4ds-lab3>

Questions?

Docker Compose

Slides available at <http://bit.ly/d4ds-compose-slides>

Putting it Together

Talk Recommender

Agenda

- Introduce problem
- Build solution in Jupyter Notebook
- Deploy solution as a service



32 Tutorials

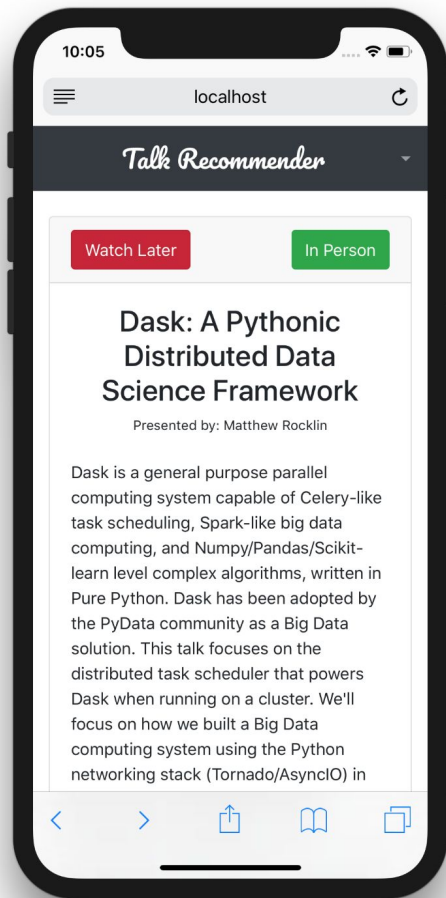
12 Sponsor Workshops

16 Education Summit Talks

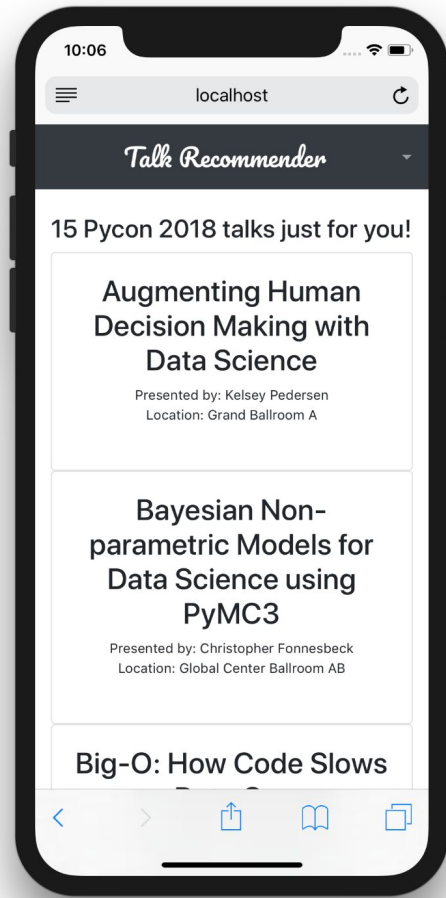
95 Conference Talks



32 Tutorials
12 Sponsor Workshops
16 Education Summit Talks
95 Conference Talks



iPhone X - 11.2



iPhone X - 11.2

Demo

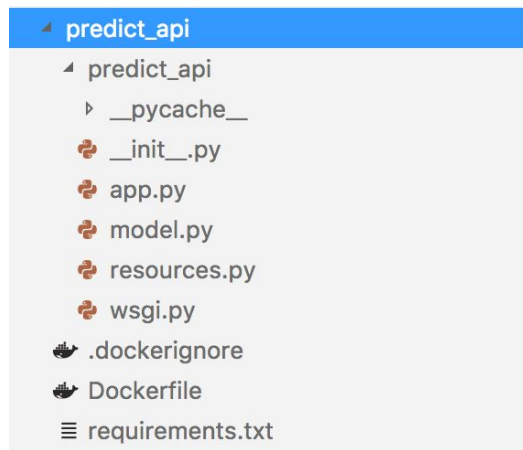
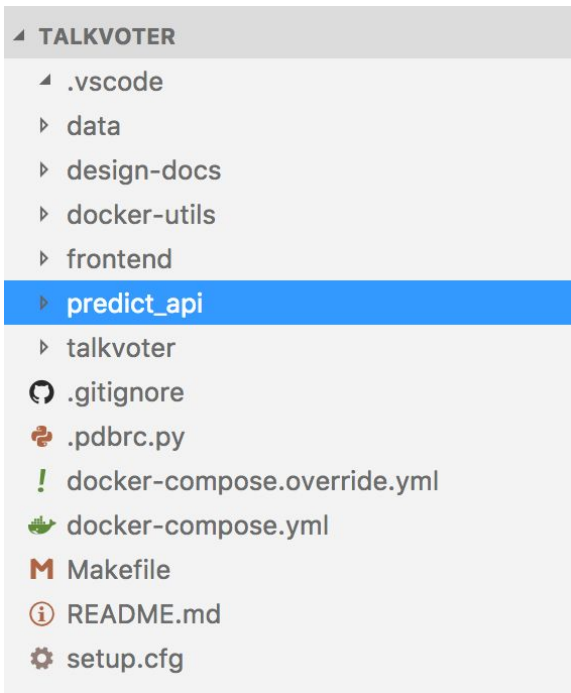
Lab: Data Science Essentials

<http://bit.ly/d4ds-lab4>

Talk Recommender: Code Walk Thru

<https://github.com/docker-for-data-science/talkvoter>

Talk Recommender: predict_api



Talk Recommender: frontend

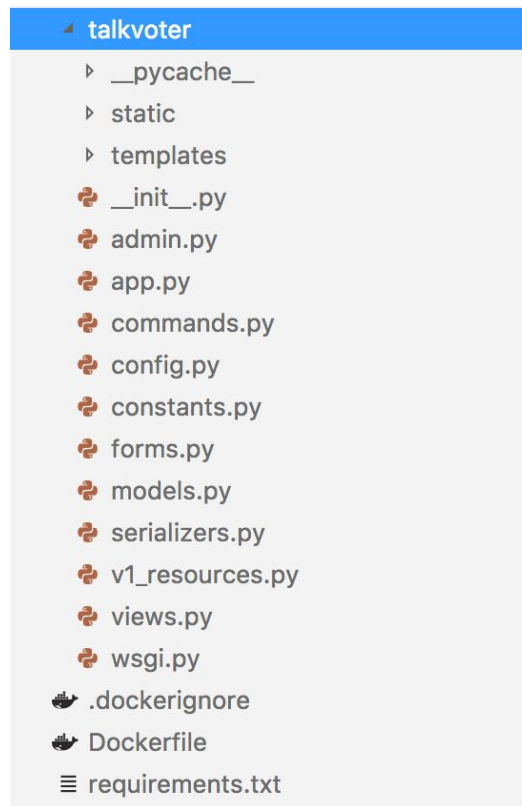
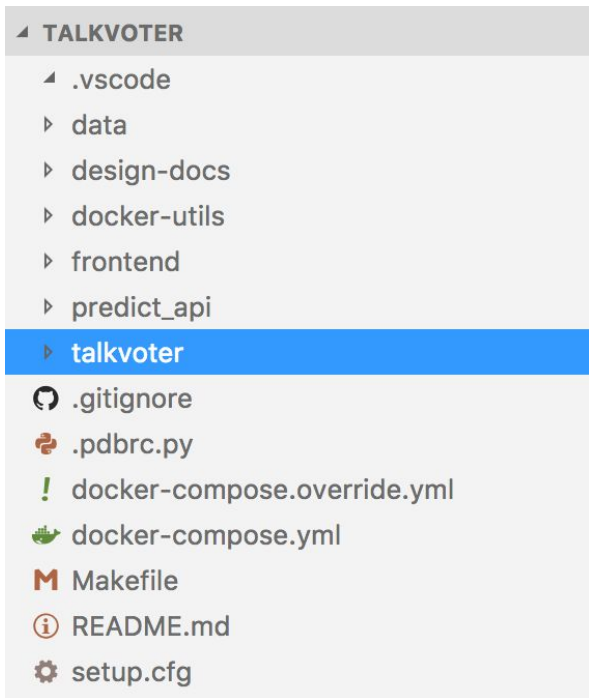
▲ TALKVOTER

- .vscode
- data
- design-docs
- docker-utils
- frontend
- predict_api
- talkvoter
- 🔗 .gitignore
- 🐍 .pdbrc.py
- ! docker-compose.override.yml
- 🐳 docker-compose.yml
- M Makefile
- 📖 README.md
- ⚙️ setup.cfg

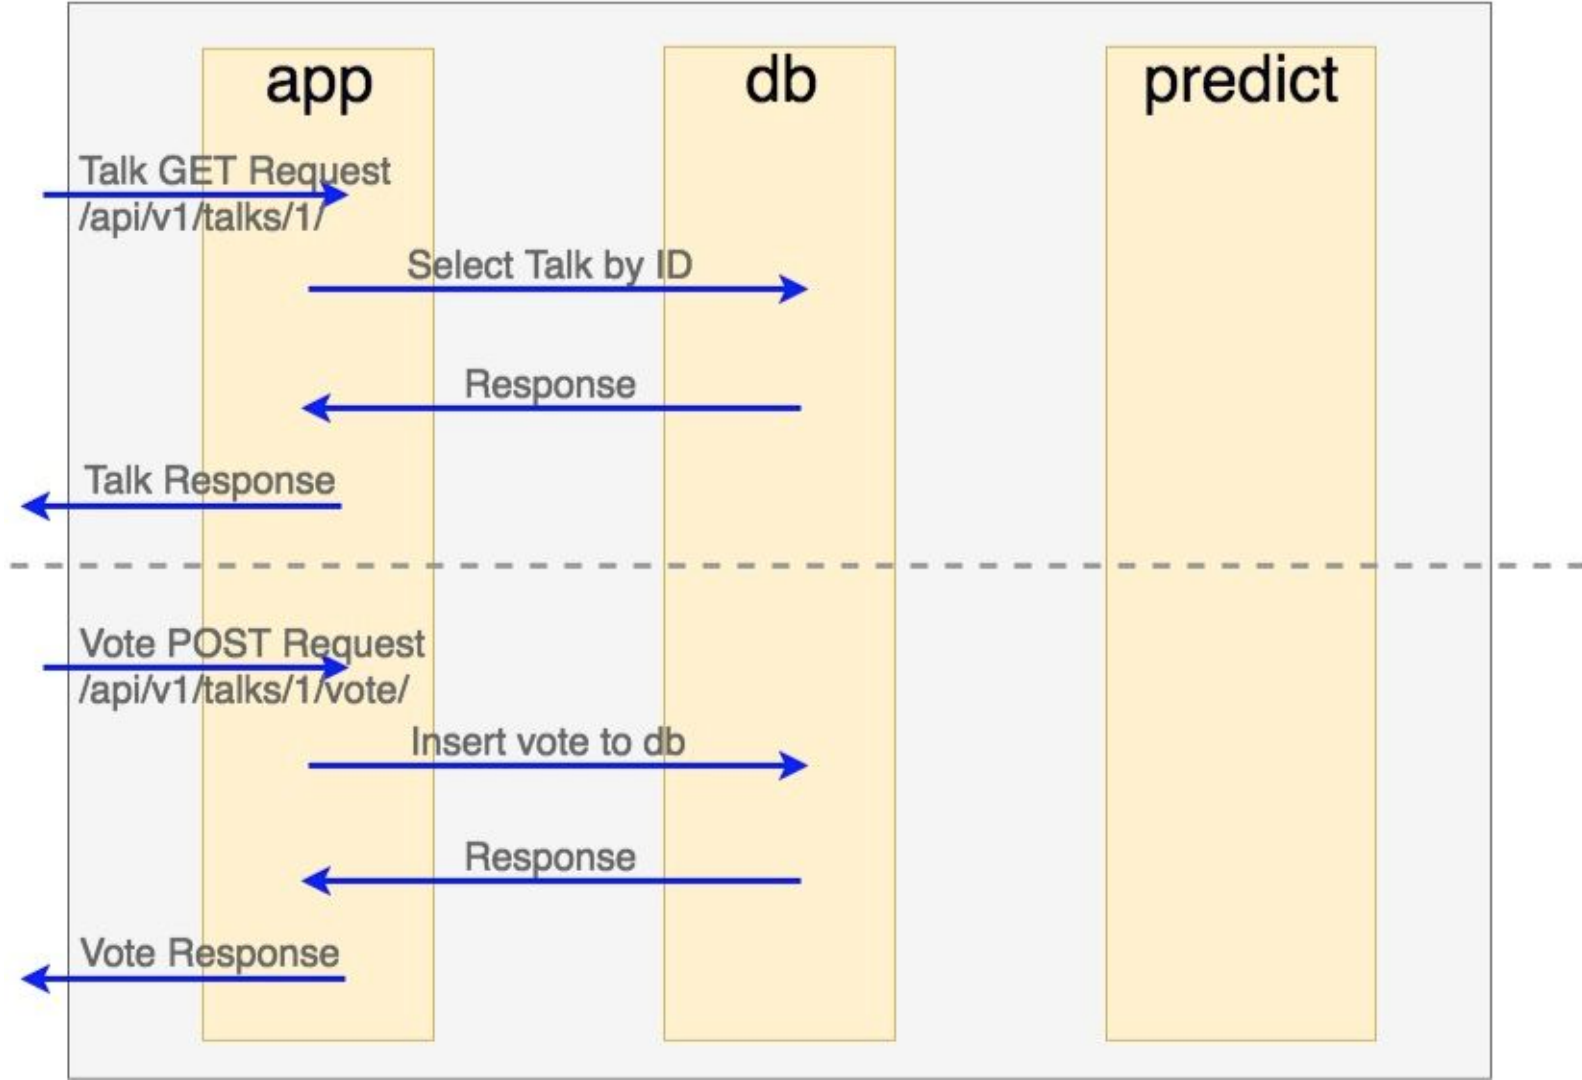
▲ frontend

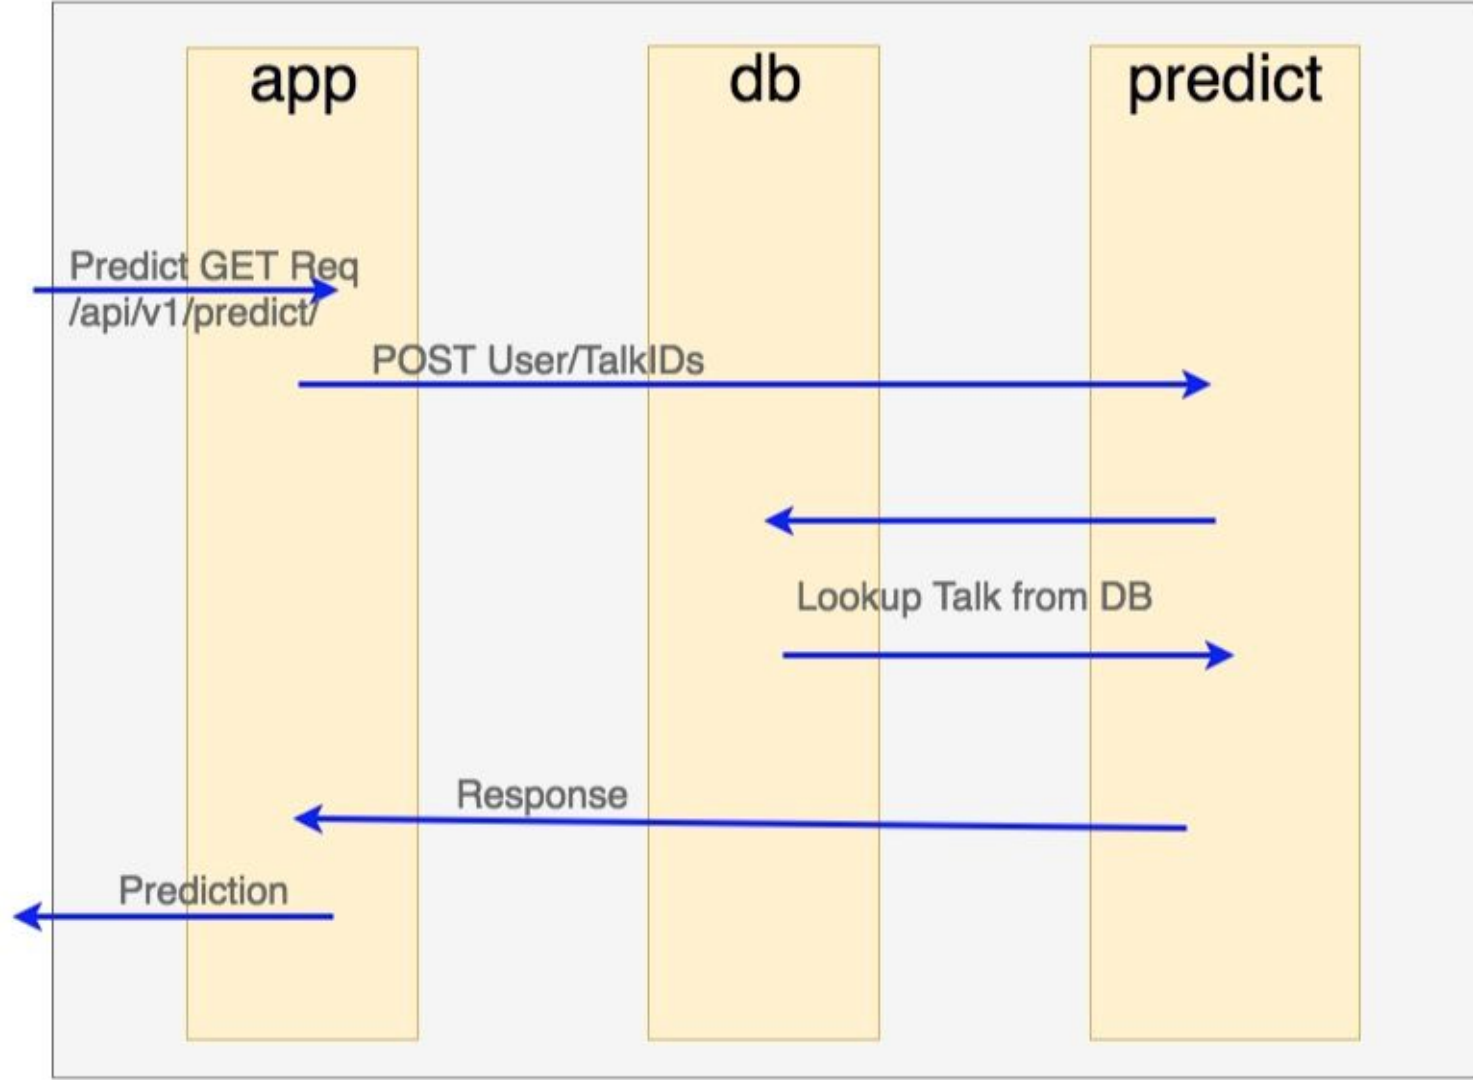
- build
- node_modules
- public
- src
- 🔗 .gitignore
- { } package-lock.json
- { } package.json
- 📖 README.md
- 👤 yarn.lock

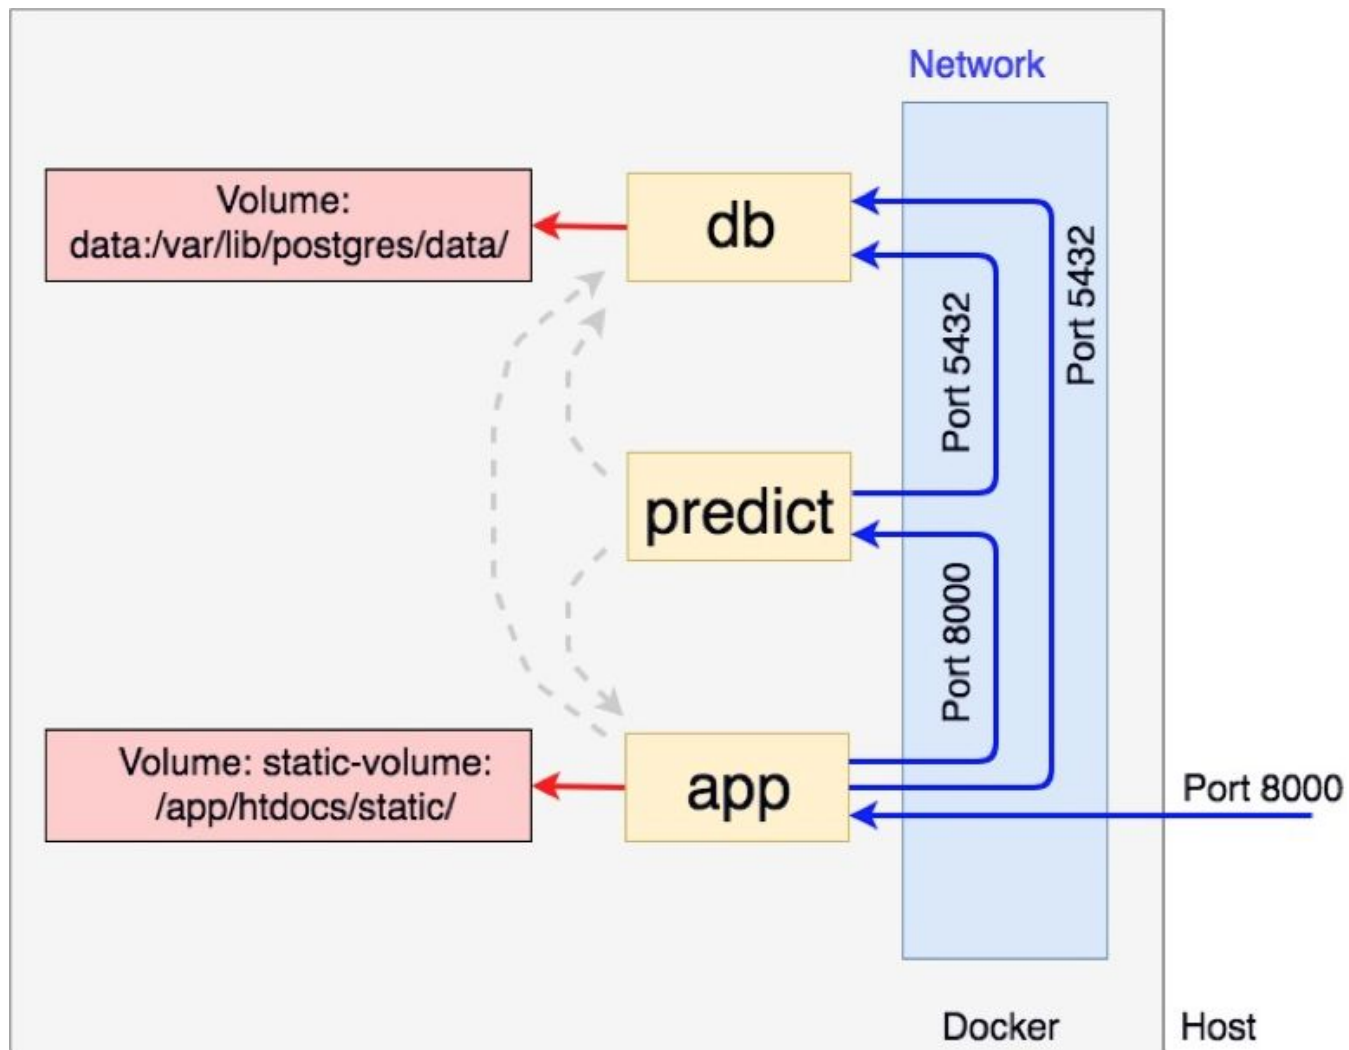
Talk Recommender: talkvoter



Talk Recommender Details







Lab: Docker-Compose Essentials



<http://bit.ly/d4ds-lab5>

Wrapping It Up

Container Workflow Best Practices

- [Use official images](#) as base when creating Dockerfile
- [Version Docker images](#), don't use `latest`
- Think of containers as [immutable objects](#)
 - To propagate changes to container, create new image
 - Use image to generate new container
- Use [multi-stage builds](#) to keep production image small
 - Copy build artifacts into final image from intermediate build image
- Do not run processes in container as [root](#)

Meet the New Excuse (Same as the Old Excuse)



Daniel Stori {turnoff.us}

Source: turnoff.us

Next Steps & Additional Resources

- [How to Install Docker](#)
- [Docker Documentation: Getting Started Guide](#)
- [Nigel Poulton's Docker Deep Dive Course](#)
- [CenturyLink Developer Center](#)
- [Kubernetes](#)

Thank You

Slides: <http://bit.ly/d4ds-slides>

Github: <http://bit.ly/d4ds-tutorial>

Twitter: [@CaiusSivjus](#) | [@JoeJJasinski](#) | [@Tathagata](#)

Acknowledgements (Easter Egg)

- [ChiPy](#)